

STATS-506 HW2

Zekai Xu

2025-09-13

[Github Repo](#)

Question 1

Part a

```
#' Compute random walk using explicit loops
#' @parameter n, a positive integer, number of steps
#' @return an integer, position after random walk
random_walk1 <- function(n)
{
  current <- 0
  for (i in 1:n)
  {
    direction <- runif(1)
    if (direction <= 0.5 ) # Left
    {
      boost <- runif(1)
      if (boost <= 0.8) # No boost
      {
        current <- current - 1
      }
      else # Boost
      {
        current <- current - 3
      }
    }
    else
```

```

{
  boost <- runif(1)
  if (boost <= 0.95) # No boost
  {
    current <- current + 1
  }
  else # Boost
  {
    current <- current + 10
  }
}
return (current)
}

```

```

#' Compute random walk using built-in R vectorized functions
#' @param n, a positive integer, number of steps
#' @return an integer, position after random walk
random_walk2 <- function(n)
{
  current <- 0

  rand_nums <- runif(2 * n)

  direction <- rand_nums[c(TRUE, FALSE)]
  direction <- as.integer(direction > 0.5)

  boost <- rand_nums[c(FALSE, TRUE)]
  boost <- as.integer((direction == 0 & boost > 0.8) | (direction == 1 & boost > 0.95))

  steps <- ifelse(direction == 1, 1 + 9 * boost, -1 - 2 * boost)
  return (sum(steps))
}

```

```

#' Compute random walk using one of the `apply` functions
#' @param n, a positive integer, number of steps
#' @return an integer, position after random walk
random_walk3 <- function(n)
{
  steps <- sapply(1:n, \i){
    direction <- runif(1)
    if (direction <= 0.5) # Left

```

```
{  
  boost <- runif(1)  
  if (boost <= 0.8)  
    return (-1)  
  else  
    return (-3)  
}  
else # Right  
{  
  boost <- runif(1)  
  if (boost <= 0.95)  
    return (1)  
  else  
    return (10)  
}  
}  
)  
return (sum(steps))  
}
```

```
random_walk1(10)
```

```
[1] -4
```

```
random_walk2(10)
```

```
[1] 4
```

```
random_walk3(10)
```

```
[1] -10
```

```
random_walk1(1000)
```

```
[1] 120
```

```
random_walk2(1000)
```

```
[1] 45
```

```
random_walk3(1000)
```

```
[1] -65
```

Part b

```
SEED <- 43
```

```
set.seed(SEED)  
random_walk1(10)
```

```
[1] -10
```

```
random_walk1(1000)
```

```
[1] 150
```

```
set.seed(SEED)  
random_walk2(10)
```

```
[1] -10
```

```
random_walk2(1000)
```

```
[1] 150
```

```
set.seed(SEED)  
random_walk3(10)
```

```
[1] -10
```

```
random_walk3(1000)
```

```
[1] 150
```

Part c

```

TIMES <- 10
N1 <- 1000
N2 <- 100000

result_1k <- microbenchmark(random_walk1(N1), random_walk2(N1),
                             random_walk3(N1), times = TIMES)

```

Warning in microbenchmark(random_walk1(N1), random_walk2(N1), random_walk3(N1),
: less accurate nanosecond times to avoid potential integer overflows

```

result_100k <- microbenchmark(random_walk1(N2), random_walk2(N2),
                               random_walk3(N2), times = TIMES)

print(result_1k)

```

	expr	min	lq	mean	median	uq	max	neval
random_walk1(N1)	730.907	755.466	792.1938	793.514	819.754	859.483	10	
random_walk2(N1)	41.492	46.699	50.9917	51.086	56.621	58.753	10	
random_walk3(N1)	993.512	1038.694	1067.7876	1071.761	1098.021	1130.493	10	

```
print(result_100k)
```

	expr	min	lq	mean	median	uq	max	neval
random_walk1(N2)	81.440227	82.138867	83.382664	82.58774	84.503829			
random_walk2(N2)	3.860519	3.923536	4.339395	4.16191	4.324926			
random_walk3(N2)	112.537538	115.168385	128.515472	118.74971	121.815592			

Based on the benchmark results, the vectorized implementation (random_walk2) is consistently the fastest, being nearly twenty times quicker than the explicit loop (random_walk1) and over twenty-five times faster than the apply-family version (random_walk3) for both small and large input sizes. The explicit loop performs moderately well, faster than the apply approach but still much slower than vectorization, while the apply family is the slowest due to the overhead of repeated function calls. Overall, vectorization clearly outperforms the other methods, especially as the problem size grows.

Part d

The random walk is $Y = \sum_{i=1}^n X_i$, where X_i i.i.d follows discrete pdf:

$$P(X_i = -3) = 0.1, \quad P(X_i = -1) = 0.4, \quad P(X_i = 1) = 0.475, \quad P(X_i = 10) = 0.025$$

Note X_i has finite variance, therefore we could apply central limit theorem:

$$\frac{\bar{X} - E[X_i]}{\sqrt{Var(X_i)/\sqrt{n}}} \xrightarrow{d} N(0, 1) \Rightarrow Y \xrightarrow{d} N(nE[X_i], nVar(X_i))$$

To compute the probability, we adopt **local** central limit theorem, and the approximate probability of $Y = 0$ can be calculated as follow:

$$P(Y = 0) \approx \frac{h}{\sqrt{2\pi n Var(X_i)}} \exp\left\{-\frac{1}{2nVar(X_i)}(0 - nE[X_i])^2\right\}$$

where h is the lattice span (the greatest common divisor of all pairwise differences between possible values of X_i), i.e. the spacing of the grid on which the distribution is supported. In this question, $h = 1$.

The expectation and variance of X_i are:

$$E[X_i] = 0.025 \\ Var(X_i) = E[X_i^2] - E^2[X_i] = 4.274375$$

Now we could compute the asymptotic probability with number of steps 10, 100 and 1000:

$$P(Y = 0; n = 10) = 0.06097562 \\ P(Y = 0; n = 100) = 0.01915573 \\ P(Y = 0; n = 1000) = 0.00567182$$

Monte Carlo Simulation

```
#' Monte Carlo Simulation of random walk back to 0 after n steps
#' @param n, a positive integer, number of steps
#' @param N, a positive integer, number of simulations
#' @param seed, a positive integer, random number seed
#' @return frequency of random walk back to 0 after n steps
simul <- function(n, N = 1e4, seed = SEED)
{
  vals <- c(-3, -1, 1, 10)
  probs <- c(0.1, 0.4, 0.475, 0.025)
```

```

set.seed(seed)
steps <- matrix(sample(vals, size = N * n, replace = TRUE, prob = probs), nrow = N)
sums <- rowSums(steps)

return (sum(sums == 0) / N)
}
simul(10)

```

[1] 0.1388

```
simul(100)
```

[1] 0.0195

```
simul(1000)
```

[1] 0.0055

Notice that for $n = 100$ and $n = 1000$, the theoretical asymptotic probability obtained from the local CLT is very close to the empirical frequency from the Monte Carlo simulation, which supports the validity of the local CLT approximation in large samples. However, when n is small (e.g., $n = 10$), the discrepancy between the theoretical probability and the simulation result is substantial, reflecting the inaccuracy of the asymptotic approximation in the small-sample regime.

Question 2

```

set.seed(SEED)
(
  sum(rpois(7, 1)) # Midnight - 7AM
+ sum(rnorm(2, 60, sqrt(12))) # 7AM - 9AM
+ sum(rpois(7, 7)) # 9AM - 4PM
+ sum(rnorm(2, 60, sqrt(12))) # 4PM - 6PM
+ sum(rpois(6, 12)) # 6PM - Midnight
) / 24

```

[1] 15.11011

Question 3

Part a

```
columns2drop <- c("brand", "superbowl_ads_dot_com_url", "youtube_url",
                  "id", "etag", "published_at", "title", "description",
                  "thumbnail", "channel_title", "kind")
youtube_deid <- youtube %>%
  select(-all_of(columns2drop))
dim(youtube_deid)
```

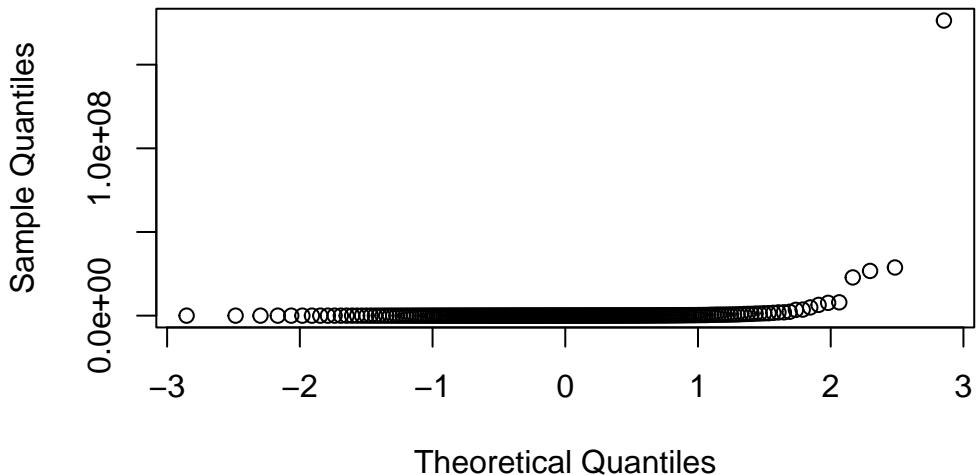
```
[1] 247 14
```

Part b

View counts: category 2

```
qqnorm(youtube_deid$view_count)
```

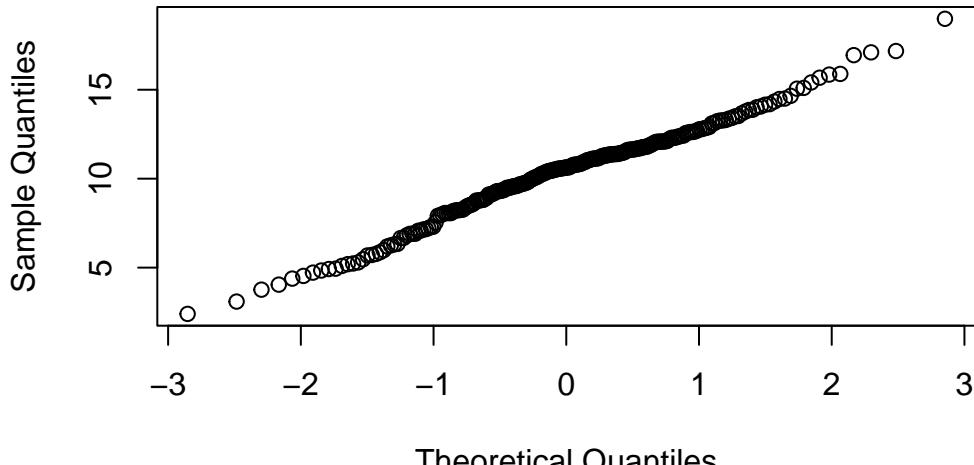
Normal Q–Q Plot



The View counts column is heavily right-skewed, and is therefore not suitable to serve as outcome variable directly

```
qqnorm(log(1 + youtube_deid$view_count))
```

Normal Q–Q Plot



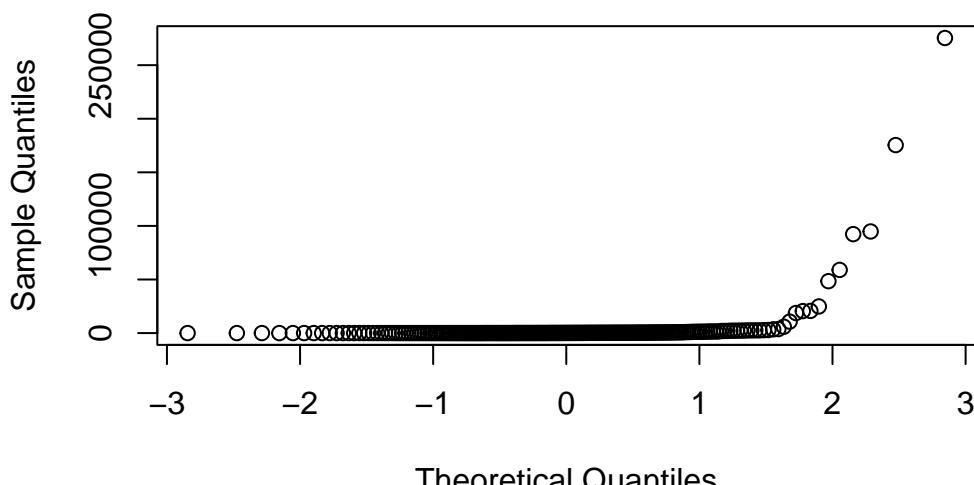
Through transformation $\log(1 + x)$, the data is approximately normal, and is therefore appropriate to serve as outcome variable

```
youtube_deid$view_count <- log(1 + youtube_deid$view_count)
```

Like counts: category 2

```
qqnorm(youtube_deid$like_count)
```

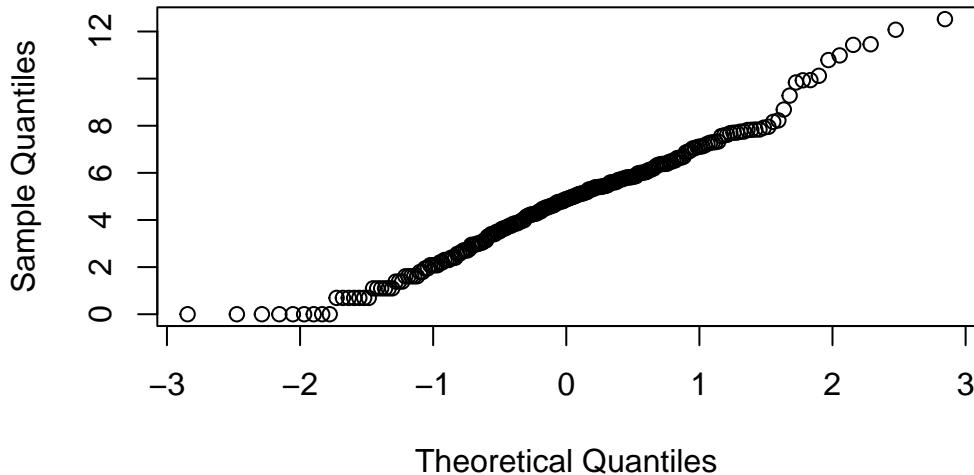
Normal Q–Q Plot



The Like counts column is heavily right-skewed, and is therefore not suitable to serve as outcome variable directly

```
qqnorm(log(1 + youtube_deid$like_count))
```

Normal Q-Q Plot




Through transformation $\log(1 + x)$, the data is approximately normal, and is therefore appropriate to serve as outcome variable

```
youtube_deid$like_count <- log(1 + youtube_deid$like_count)
```

Dislike counts: category 2

```
qqnorm(youtube_deid$dislike_count)
```

Normal Q–Q Plot




Theoretical Quantiles

The Dislike counts column is heavily right-skewed, and is therefore not suitable to serve as outcome variable directly

```
qqnorm(log(1 + youtube_deid$dislike_count))
```

Normal Q–Q Plot



Theoretical Quantiles

Through transformation $\log(1 + x)$, the data is approximately normal, and is therefore appropriate to serve as outcome variable

```
youtube_deid$dislike_count <- log(1 + youtube_deid$dislike_count)
```

Favorite counts: category 3

```
unique(youtube_deid$favorite_count)
```

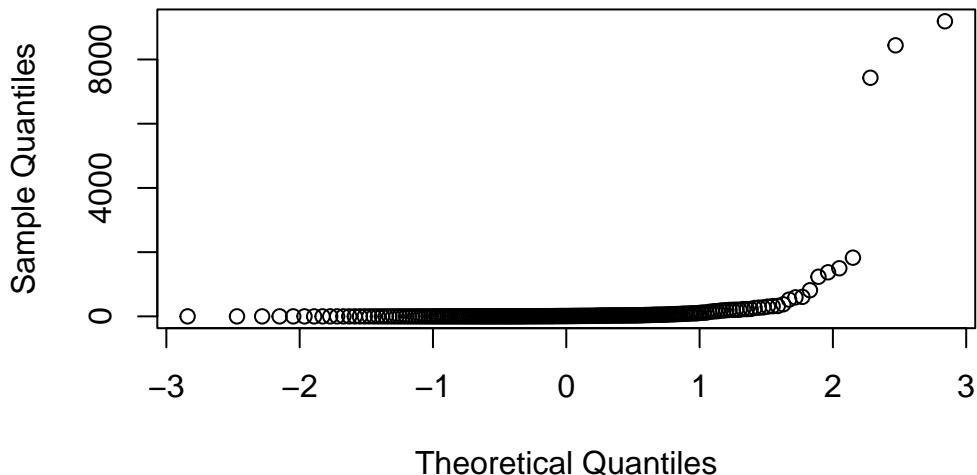
```
[1] 0 NA
```

The Favorite counts only has 0 and NA values, which is categorical, and this column is therefore not suitable to serve as outcome variable.

Comment counts: category 2

```
qqnorm(youtube_deid$comment_count)
```

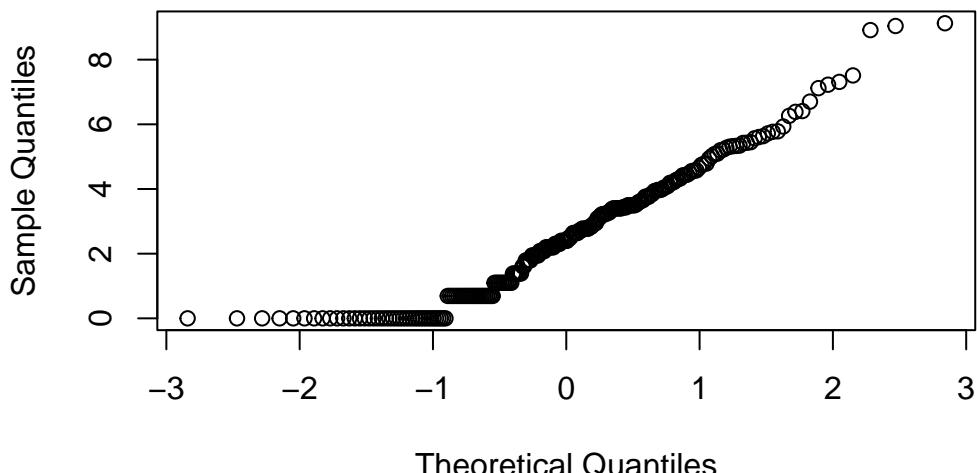
Normal Q–Q Plot



The Comment counts column is heavily right-skewed, and is therefore not suitable to serve as outcome variable directly

```
qqnorm(log(1 + youtube_deid$comment_count))
```

Normal Q-Q Plot



Through transformation $\log(1 + x)$, the data is approximately normal, and is therefore appropriate to serve as outcome variable

```
youtube_deid$comment_count <- log(1 + youtube_deid$comment_count)
```

Part c

```
outcomes <- c("view_count", "like_count", "dislike_count", "comment_count")
predictors <- "funny + show_product_quickly + patriotic +
               celebrity + danger + animals + use_sex + year"
models <- lapply(outcomes, function(y)
{
  formula <- as.formula(paste(y, "~", predictors))
  lm(formula, data = youtube_deid)
})

lapply(models, summary)
```

```
[[1]]
```

```
Call:
lm(formula = formula, data = youtube_deid)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-7.7742	-1.6152	0.1311	1.7036	8.4481

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-31.55016	71.00538	-0.444	0.657
funnyTRUE	0.56492	0.46702	1.210	0.228
show_product_quicklyTRUE	0.21089	0.40530	0.520	0.603
patrioticTRUE	0.50699	0.53811	0.942	0.347
celebrityTRUE	0.03548	0.42228	0.084	0.933
dangerTRUE	0.63131	0.41812	1.510	0.132
animalsTRUE	-0.31002	0.39348	-0.788	0.432
use_sexTRUE	-0.38671	0.44782	-0.864	0.389
year	0.02053	0.03531	0.582	0.561

Residual standard error: 2.787 on 222 degrees of freedom

(16 observations deleted due to missingness)

Multiple R-squared: 0.02694, Adjusted R-squared: -0.008122
F-statistic: 0.7684 on 8 and 222 DF, p-value: 0.631

[[2]]

Call:

lm(formula = formula, data = youtube_deid)

Residuals:

Min	1Q	Median	3Q	Max
-5.2860	-1.6333	0.0868	1.4911	7.7431

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-150.51357	63.45723	-2.372	0.0186 *
funnyTRUE	0.47476	0.41816	1.135	0.2575
show_product_quicklyTRUE	0.20017	0.36391	0.550	0.5828
patrioticTRUE	0.80689	0.49791	1.621	0.1066
celebrityTRUE	0.41283	0.38212	1.080	0.2812
dangerTRUE	0.63895	0.37350	1.711	0.0886 .
animalsTRUE	-0.05944	0.35298	-0.168	0.8664
use_sexTRUE	-0.42952	0.40064	-1.072	0.2849
year	0.07685	0.03155	2.436	0.0157 *
<hr/>				
Signif. codes:	0 ***	0.001 **	0.01 *'	0.05 '.' 0.1 ' ' 1

```
Residual standard error: 2.467 on 216 degrees of freedom  
(22 observations deleted due to missingness)  
Multiple R-squared:  0.07313,   Adjusted R-squared:  0.03881  
F-statistic:  2.13 on 8 and 216 DF,  p-value: 0.0342
```

```
[[3]]
```

```
Call:  
lm(formula = formula, data = youtube_deid)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-4.0292	-1.3299	-0.3192	0.8986	8.7219

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-183.06813	53.34768	-3.432	0.000719 ***
funnyTRUE	0.25949	0.35154	0.738	0.461224
show_product_quicklyTRUE	0.27511	0.30593	0.899	0.369515
patrioticTRUE	0.81407	0.41859	1.945	0.053095 .
celebrityTRUE	-0.20214	0.32125	-0.629	0.529852
dangerTRUE	0.22184	0.31400	0.707	0.480630
animalsTRUE	-0.21192	0.29675	-0.714	0.475911
use_sexTRUE	-0.32980	0.33681	-0.979	0.328583
year	0.09207	0.02653	3.471	0.000626 ***

Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	'	'	'	'
	'	'	'	'

```
Residual standard error: 2.074 on 216 degrees of freedom  
(22 observations deleted due to missingness)  
Multiple R-squared:  0.09753,   Adjusted R-squared:  0.06411  
F-statistic: 2.918 on 8 and 216 DF,  p-value: 0.004115
```

```
[[4]]
```

```
Call:  
lm(formula = formula, data = youtube_deid)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

```

-4.1372 -1.4665 -0.1427  1.4061  5.8468

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) -99.09835  52.92351 -1.872  0.0625 .
funnyTRUE      0.21954   0.34528   0.636  0.5256
show_product_quicklyTRUE 0.40939   0.30229   1.354  0.1771
patrioticTRUE    0.66698   0.39902   1.672  0.0961 .
celebrityTRUE     0.29767   0.31541   0.944  0.3464
dangerTRUE        0.17784   0.31069   0.572  0.5677
animalsTRUE       -0.26802   0.29347  -0.913  0.3621
use_sexTRUE       -0.39323   0.33163  -1.186  0.2370
year              0.05034   0.02632   1.913  0.0571 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.039 on 213 degrees of freedom
(25 observations deleted due to missingness)
Multiple R-squared:  0.06535, Adjusted R-squared:  0.03025
F-statistic: 1.862 on 8 and 213 DF,  p-value: 0.06748

```

Across the four linear regression models, the advertising characteristics showed little consistent association with YouTube engagement metrics, as most binary flags were statistically insignificant. The only notable effects were that ads featuring danger tended to receive more likes, and patriotic ads showed some evidence of attracting more dislikes and comments, though these results were only marginally significant. In contrast, year was a consistent predictor: newer ads were associated with significantly higher numbers of likes and dislikes, and a marginally higher number of comments, reflecting the overall growth of YouTube engagement over time. However, the explanatory power of all models was low (R^2 values below 0.1), indicating that engagement is likely driven by other unobserved factors such as brand influence, ad content quality, or promotion strategy.

Part d

```

vars <- c("view_count", "funny", "show_product_quickly", "patriotic",
        "celebrity", "danger", "animals", "use_sex", "year")
# Filter na
dat <- youtube_deid[, vars]
dat <- dat[complete.cases(dat),]

```

```

form <- ~ funny + show_product_quickly + patriotic + celebrity + danger + animals + use_sex +
X <- model.matrix(form, data = dat)
y <- dat$view_count

XtX    <- crossprod(X)
XtX_i <- solve(XtX)
Xty    <- crossprod(X, y)
beta_hat <- XtX_i %*% Xty

beta_hat

```

	[,1]
(Intercept)	-31.55015804
funnyTRUE	0.56492445
show_product_quicklyTRUE	0.21088918
patrioticTRUE	0.50699051
celebrityTRUE	0.03547862
dangerTRUE	0.63131085
animalsTRUE	-0.31001838
use_sexTRUE	-0.38670726
year	0.02053399

```
summary(models[[1]])
```

Call:
`lm(formula = formula, data = youtube_deid)`

Residuals:

Min	1Q	Median	3Q	Max
-7.7742	-1.6152	0.1311	1.7036	8.4481

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-31.55016	71.00538	-0.444	0.657
funnyTRUE	0.56492	0.46702	1.210	0.228
show_product_quicklyTRUE	0.21089	0.40530	0.520	0.603
patrioticTRUE	0.50699	0.53811	0.942	0.347
celebrityTRUE	0.03548	0.42228	0.084	0.933
dangerTRUE	0.63131	0.41812	1.510	0.132
animalsTRUE	-0.31002	0.39348	-0.788	0.432

use_sexTRUE	-0.38671	0.44782	-0.864	0.389
year	0.02053	0.03531	0.582	0.561

Residual standard error: 2.787 on 222 degrees of freedom

(16 observations deleted due to missingness)

Multiple R-squared: 0.02694, Adjusted R-squared: -0.008122

F-statistic: 0.7684 on 8 and 222 DF, p-value: 0.631