

# STATS-506 HW3

Zekai Xu

2025-09-29

[Github Repo](#)

## Question 1

### Part a

```
# Load Dataset
aux <- as.data.frame(read_xpt("./AUX_I.xpt"))
demo <- as.data.frame(read_xpt("./DEMO_I.xpt"))

# Merge Dataset
df <- inner_join(aux, demo, by = "SEQN")

# Print dimension of the merged dataframe
print(dim(df))
```

[1] 4582 119

### Part b

```
# Select columns used in this question
var2keep <- c(
  "SEQN", #Identifier
  "RIAGENDR", #Gender
  "DMDCITZN", #Citizenship
  "DMDHHSZA", #Number of children 5 years or younger in the household
```

```

"INDHHIN2", #Annual household income
"AUXTWIDR", #Tympanometric width, right ear
"AUXTWIDL" #Tympanometric width, left ear
)
df_partb <- df[var2keep]

# Rename columns
names(df_partb) <- c("id", "gender", "citizenship", "numOfKids", "income", "tw_re", "tw_le")

# Convert "unknown" data points to NA values
df_partb$citizenship[df_partb$citizenship %in% c(77, 99)] <- NA
df_partb$income[df_partb$income %in% c(77, 99)] <- NA

# Convert categorical column into factor
df_partb$gender <- factor(
  df_partb$gender,
  levels = c(1, 2),
  labels = c("Male", "Female")
)
df_partb$citizenship <- factor(
  df_partb$citizenship,
  levels = c(1, 2),
  labels = c("USA", "Other")
)
#df_partb$numOfKids <- factor(df_partb$numOfKids)
#df_partb$income <- factor(df_partb$income)

# Overview of processed data
glimpse(df_partb)

```

Rows: 4,582

Columns: 7

\$ id	<dbl>	83732, 83733, 83735, 83736, 83741, 83742, 83744, 83747, 83~
\$ gender	<fct>	Male, Male, Female, Female, Male, Female, Male, Male, Male~
\$ citizenship	<fct>	USA, Other, USA, USA, USA, Other, USA, USA, USA, USA, USA, USA,~
\$ numOfKids	<dbl>	0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0~
\$ income	<dbl>	10, 4, 10, 7, 7, 6, 3, 3, 10, 15, 6, 5, 5, 1, 8, 14, 10, 2~
\$ tw_re	<dbl>	49, 83, 110, 99, 73, 99, 122, 69, 63, 78, 44, 133, 144, NA~
\$ tw_le	<dbl>	76, 51, 49, 116, 73, 70, 111, 73, 60, 88, 39, 112, NA, NA,~

## Part c

```
# Create formulas for 4 Poisson Regression model
preds_1r <- c("gender")
preds_2r <- c("gender", "citizenship", "numOfKids", "income")
preds_1l <- c("gender")
preds_2l <- c("gender", "citizenship", "numOfKids", "income")
fm_1r <- reformulate(term.labels = preds_1r, response = "tw_re", intercept = TRUE)
fm_2r <- reformulate(term.labels = preds_2r, response = "tw_re", intercept = TRUE)
fm_1l <- reformulate(term.labels = preds_1l, response = "tw_le", intercept = TRUE)
fm_2l <- reformulate(term.labels = preds_2l, response = "tw_le", intercept = TRUE)

# Model
m_1r <- glm(formula = fm_1r, data = df_partb, family = poisson(link = "log"))
m_2r <- glm(formula = fm_2r, data = df_partb, family = poisson(link = "log"))
m_1l <- glm(formula = fm_1l, data = df_partb, family = poisson(link = "log"))
m_2l <- glm(formula = fm_2l, data = df_partb, family = poisson(link = "log"))
```

Table 1: Incidence Rate Ratios (IRR) with 95 percent CI and p-values

Model	Term	Effect Size	
		IRR [95 percent CI]	p-value
1R (Right: gender)	genderFemale	1.01 [1.00, 1.02]	0.006
2R (Right: gender + citizenship + numOfKids + income)	genderFemale	1.01 [1.01, 1.02]	<0.001
2R (Right: gender + citizenship + numOfKids + income)	citizenshipOther	1.07 [1.07, 1.08]	<0.001
2R (Right: gender + citizenship + numOfKids + income)	numOfKids	1.00 [1.00, 1.01]	0.860
2R (Right: gender + citizenship + numOfKids + income)	income	0.99 [0.99, 1.00]	<0.001
1L (Left: gender)	genderFemale	1.01 [1.01, 1.02]	<0.001
2L (Left: gender + citizenship + numOfKids + income)	genderFemale	1.02 [1.01, 1.02]	<0.001
2L (Left: gender + citizenship + numOfKids + income)	citizenshipOther	1.04 [1.03, 1.05]	<0.001
2L (Left: gender + citizenship + numOfKids + income)	numOfKids	0.98 [0.98, 0.99]	<0.001
2L (Left: gender + citizenship + numOfKids + income)	income	1.00 [1.00, 1.00]	<0.001

## Part d

```
# ----- (A) Likelihood Ratio Test: does gender improve fit? -----
m_21_nogender <- update(m_21, . ~ . - gender)
anova(m_21_nogender, m_21, test = "LRT") # reports Chi-sq, df, p-value
```

Table 2: Model Fit Statistics: Sample Size, McFadden Pseudo-R2, and AIC

Model	N	Pseudo-R2 (McFadden)	AIC
1R	4,149	0.000	96,618
2R	3,886	0.067	90,169
1L	4,103	0.000	98,685
2L	3,847	0.070	91,835

#### Analysis of Deviance Table

```

Model 1: tw_le ~ citizenship + numOfKids + income
Model 2: tw_le ~ gender + citizenship + numOfKids + income
  Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1       3843     68040
2       3842     68021  1   19.318 1.107e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

# ----- (B) IRR difference (Wald/contrast): Female vs Male -----
# emmeans on link scale, then contrast on response scale -> ratio = IRR
emm <- emmeans(m_21, ~ gender)           # marginal means (on link scale)
pairs(emm) %>% summary(type = "response") # ratio, 95% CI, p-value (IRR Female/Male)

```

```

contrast      ratio      SE  df null z.ratio p.value
Male / Female 0.985 0.00346 Inf    1  -4.394  <.0001

```

Results are averaged over the levels of: citizenship  
 Tests are performed on the log scale

```

# ----- (C) Predicted mean TW by gender (adjusted) -----
# Response-scale marginal means (i.e., predicted means for each gender)
summary(emm, type = "response")          # mean, SE, 95% CI for each gender

```

```

gender rate      SE  df asympt.LCL asympt.UCL
Male   85.7 0.246 Inf      85.2      86.1
Female 87.0 0.243 Inf      86.5      87.5

```

Results are averaged over the levels of: citizenship  
 Confidence level used: 0.95  
 Intervals are back-transformed from the log scale

## Evidence of Gender Difference (Model 2L)

**Task.** Assess whether males and females differ in (i) their **incidence rate ratio (IRR)** and (ii) the **predicted value** of left-ear tympanometric width (TW), adjusting for citizenship, number of children, and income.

### 1) Likelihood-Ratio Test (model-level evidence)

- Compare Model 1 (without gender) vs Model 2L (with gender).
- $\Delta\text{Deviance} = \mathbf{19.318}$  on 1 df,  $p = \mathbf{1.11 \times 10^{-4}}$ .
- **Conclusion:** Reject  $H: \text{no gender effect}$ . Adding gender significantly improves model fit.

### 2) IRR Difference (contrast/Wald test)

- **IRR (Male / Female) = 0.985**,  $z = -4.394$ ,  $p < 0.0001$ .
- Equivalently, **IRR (Female / Male) = 1.015** → females have about **1.5% higher** expected left-ear TW than males, holding covariates fixed.
- Interpretation: IRR differs significantly from 1, indicating a statistically significant gender effect on the rate (mean) of TW.

### 3) Adjusted Predicted Means (response scale)

- **Male:** 85.7 (95% CI: **85.2–86.1**)
- **Female:** 87.0 (95% CI: **86.5–87.5**)
- **Conclusion:** Predicted means differ in the same direction as the IRR result; females have a slightly higher adjusted TW.

## Overall Interpretation

Both the model-level LRT and the IRR/mean comparisons show a **statistically significant** gender difference in left-ear TW after adjustment. The **effect size is small** ( 1.5% higher for females; **1.3 units** difference in predicted means), so practical significance should be considered alongside statistical significance.

## Question 2

### Part a

```
# Connect Database
file_path <- "./sakila_master.db"
con <- dbConnect(SQLite(), dbname = file_path)

## ===== Method 1: extract -> compute in R =====
method1 <- function() {
  stores <- dbGetQuery(con, "SELECT store_id FROM store")

  customers <- dbGetQuery(con, "
    SELECT store_id, active
    FROM customer
  ") %>%
    mutate(                         # normalize 'active' to 0/1
      active = as.integer(tolower(trimws(as.character(active)))) %in%
        c("1","t","true","y","yes"))
  )

  per_store <- customers %>%
    group_by(store_id) %>%
    summarise(
      n_customers = n(),
      n_active     = sum(active),
      .groups = "drop"
    ) %>%
    mutate(pct_active = ifelse(n_customers > 0, 100 * n_active / n_customers, NA_real_))

  stores %>%
    left_join(per_store, by = "store_id") %>%
    select(store_id, n_customers, pct_active) %>%
    arrange(store_id)
}

## ===== Method 2: single SQL =====
method2 <- function() {
  dbGetQuery(con, "
    SELECT
      s.store_id,
```

```

        COUNT(c.customer_id) AS n_customers,
        CASE
            WHEN COUNT(c.customer_id) = 0 THEN NULL
            ELSE 100.0 * SUM(
                CASE
                    WHEN lower(c.active) IN ('1','t','true','y','yes') THEN 1
                    WHEN lower(c.active) IN ('0','f','false','n','no','') THEN 0
                    WHEN c.active = 1 THEN 1
                    WHEN c.active = 0 THEN 0
                    ELSE 0
                END
            ) * 1.0 / COUNT(c.customer_id)
        END AS pct_active
    FROM store AS s
    LEFT JOIN customer AS c
        ON c.store_id = s.store_id
    GROUP BY s.store_id
    ORDER BY s.store_id
  ")
}

## ===== Run & print =====
res_r   <- method1() %>% mutate(pct_active = round(pct_active, 2))
res_sql <- method2()  %>% mutate(pct_active = round(pct_active, 2))

print(res_r)

```

	store_id	n_customers	pct_active
1	1	326	97.55
2	2	273	97.44

```
print(res_sql)
```

	store_id	n_customers	pct_active
1	1	326	97.55
2	2	273	97.44

```

## ===== Microbenchmark =====
microbenchmark(
  two_step_R = method1(),

```

```

    single_SQL = method2(),
    times = 20L
)

```

Warning in microbenchmark(two\_step\_R = method1(), single\_SQL = method2(), :  
less accurate nanosecond times to avoid potential integer overflows

expr	min	lq	mean	median	uq	max	neval
two_step_R	2795.708	2846.896	2966.491	2942.365	2979.9620	3716.568	20
single_SQL	212.503	224.147	2819.402	231.035	246.4305	51309.778	20

Directly using SQL query is way faster than SQL + R approach.

## Part b

```

## ===== Method 1: Extract tables, then join in R =====
method1 <- function() {
  staff   <- dbGetQuery(con, "SELECT staff_id, first_name, last_name, address_id FROM staff")
  address <- dbGetQuery(con, "SELECT address_id, city_id FROM address")
  city    <- dbGetQuery(con, "SELECT city_id, country_id FROM city")
  country <- dbGetQuery(con, "SELECT country_id, country FROM country")

  staff %>%
    left_join(address, by = "address_id") %>%
    left_join(city,   by = "city_id") %>%
    left_join(country, by = "country_id") %>%
    mutate(full_name = paste(first_name, last_name)) %>%
    select(staff_id, first_name, last_name, full_name, country) %>%
    arrange(staff_id)
}

## ===== Method 2: Single SQL query (one-shot) =====
method2 <- function() {
  dbGetQuery(con, "
  SELECT
    s.staff_id,
    s.first_name,
    s.last_name,

```

```

        (s.first_name || ' ' || s.last_name) AS full_name,
        co.country
    FROM staff AS s
    JOIN address AS a ON s.address_id = a.address_id
    JOIN city    AS ci ON a.city_id   = ci.city_id
    JOIN country AS co ON ci.country_id = co.country_id
    ORDER BY s.staff_id
    ")
}

## ===== Run both & show a sample =====
res_r  <- method1()
res_sql <- method2()

print(head(res_r, 10))

  staff_id first_name last_name      full_name      country
1          1       Mike    Hillyer Mike Hillyer      Canada
2          2       Jon    Stephens Jon Stephens Australia

print(head(res_sql, 10))

  staff_id first_name last_name      full_name      country
1          1       Mike    Hillyer Mike Hillyer      Canada
2          2       Jon    Stephens Jon Stephens Australia

## ===== Microbenchmark =====
mb <- microbenchmark(
  two_step_R = method1(),
  single_SQL = method2(),
  times = 50L
)
print(mb)

Unit: microseconds
      expr      min       lq     mean   median      uq     max neval
two_step_R 2632.897 2685.869 2954.3280 2780.251 2941.791 6246.678     50
single_SQL  96.678 106.149 127.9561 113.283 120.868  777.565     50

```

Directly using SQL query is way faster than SQL + R approach.

## Part c

```
# ----- Method 1: extract tables -> compute in R -----
method1 <- function() {
  payment <- dbGetQuery(con, "SELECT rental_id, amount FROM payment")
  rental   <- dbGetQuery(con, "SELECT rental_id, inventory_id FROM rental")
  inventory<- dbGetQuery(con, "SELECT inventory_id, film_id FROM inventory")
  film     <- dbGetQuery(con, "SELECT film_id, title FROM film")

  # total $ per rental, then map rental -> film -> title
  per_rental <- payment %>%
    group_by(rental_id) %>%
    summarise(rental_value = sum(amount), .groups = "drop") %>%
    inner_join(rental, by = "rental_id") %>%
    inner_join(inventory, by = "inventory_id") %>%
    inner_join(film, by = "film_id")

  max_val <- max(per_rental$rental_value, na.rm = TRUE)

  per_rental %>%
    filter(rental_value == max_val) %>%
    distinct(title, rental_value) %>%
    arrange(title)
}

# ----- Method 2: single SQL query -----
method2 <- function() {
  dbGetQuery(con, "
    WITH per_rental AS (
      SELECT rental_id, SUM(amount) AS rental_value
      FROM payment
      GROUP BY rental_id
    ),
    max_rental AS (
      SELECT MAX(rental_value) AS max_val FROM per_rental
    )
    SELECT f.title, pr.rental_value
    FROM per_rental pr
    JOIN rental r      ON pr.rental_id = r.rental_id
    JOIN inventory i  ON r.inventory_id = i.inventory_id
    JOIN film f        ON i.film_id = f.film_id
    WHERE pr.rental_value = (SELECT max_val FROM max_rental)
  ")
}
```

```

        ORDER BY f.title
    ")
}

# ----- Run & show -----
res_r   <- method1()
res_sql <- method2()

print(res_r)

```

```

# A tibble: 9 x 2
  title          rental_value
  <chr>           <dbl>
1 FLINTSTONES HAPPINESS      12.0
2 MIDSUMMER GROUNDHOG       12.0
3 MINE TITANS              12.0
4 SCORPION APOLLO          12.0
5 SHOW LORD                 12.0
6 STING PERSONAL            12.0
7 TIES HUNGER               12.0
8 TRAP GUYS                  12.0
9 VIRTUAL SPOILERS          12.0

```

```

print(res_sql)

```

	title	rental_value
1	FLINTSTONES HAPPINESS	11.99
2	MIDSUMMER GROUNDHOG	11.99
3	MINE TITANS	11.99
4	SCORPION APOLLO	11.99
5	SCORPION APOLLO	11.99
6	SHOW LORD	11.99
7	STING PERSONAL	11.99
8	TIES HUNGER	11.99
9	TRAP GUYS	11.99
10	VIRTUAL SPOILERS	11.99

```

# ----- Microbenchmark -----
microbenchmark(
  two_step_R = method1(),

```

```
    single_SQL = method2(),
    times = 20L
)
```

```
Unit: milliseconds
      expr      min       lq     mean   median      uq     max neval
two_step_R 24.879169 25.24604 26.326065 25.761981 27.76930 28.457116    20
single_SQL  5.394903  5.46163  5.868439  6.002052  6.15859  6.744705    20
```

Directly using SQL query is way faster than SQL + R approach.

## Question 3

### Part a

```
# Load Dataset
df <- read.csv("./au-500.csv", sep = ",")  
  
# Calculate percentage
pct <- sum(str_detect(df$web, regex("\\.com$"))) / dim(df)[1] * 100
cat(sprintf("Percentage of websites that end with `\\.com`: %.2f%%\n", pct))
```

```
Percentage of websites that end with `\\.com`: 0.00%
```

### Part b

```
dom <- str_extract(df$email, regex("(?<=@)(?:[A-Za-z0-9-]+\\\.)+[A-Za-z]{2,}"))
sort(table(dom), decreasing = TRUE)[1]
```

```
hotmail.com
114
```

### Part c

```

p_excl_comma_ws <- mean(str_detect(df$company_name, "[^\\p{L},\\s]"), na.rm = TRUE)
p_excl_comma_ws_amp <- mean(str_detect(df$company_name, "[^\\p{L},\\s&]"), na.rm = TRUE)

cat(sprintf("Proportion with non-alphabetic (excluding comma/whitespace): %.2f%%\n",
            100 * p_excl_comma_ws))

```

Proportion with non-alphabetic (excluding comma/whitespace): 9.00%

```

cat(sprintf("Proportion with non-alphabetic (excluding comma/whitespace/&): %.2f%%\n",
            100 * p_excl_comma_ws_amp))

```

Proportion with non-alphabetic (excluding comma/whitespace/&): 0.80%

## Part d

```

phone_cols <- names(df)[str_detect(tolower(names(df)), "(phone|mobile)")]
fmt_au_mobile <- function(x) {
  d <- str_replace_all(x, "[^0-9]", "") # strip non-digits
  ifelse(nchar(d) == 10,
         paste0(substr(d, 1, 4), "-", substr(d, 5, 7), "-", substr(d, 8, 10)),
         NA_character_) # not 10 digits -> NA
}

cat("Before:\n")

```

Before:

```

print(utils::head(df[phone_cols], 10))

```

	phone1	phone2
1	03-8174-9123	0458-665-290
2	07-9997-3366	0497-622-620
3	08-5558-9019	0427-885-282
4	02-6044-4682	0443-795-912
5	02-1455-6085	0453-666-885

```
6 08-7868-1355 0451-966-921
7 08-6522-8931 0427-991-688
8 02-5226-9402 0415-961-606
9 07-3184-9989 0411-732-965
10 08-6890-4661 0461-862-457
```

```
df <- df %>% mutate(across(all_of(phone_cols), fmt_au_mobile))
cat("\nAfter (cell format 1234-567-890):\n")
```

After (cell format 1234-567-890):

```
print(head(df[phone_cols], 10))
```

```
  phone1      phone2
1 0381-749-123 0458-665-290
2 0799-973-366 0497-622-620
3 0855-589-019 0427-885-282
4 0260-444-682 0443-795-912
5 0214-556-085 0453-666-885
6 0878-681-355 0451-966-921
7 0865-228-931 0427-991-688
8 0252-269-402 0415-961-606
9 0731-849-989 0411-732-965
10 0868-904-661 0461-862-457
```

## Part e

```
addr_col <- names(df)[str_detect(tolower(names(df)), "address")][1]

# Extract trailing digits as apartment number; keep positive integers only
apt_num <- df[[addr_col]] %>%
  as.character() %>%
  str_extract("\\d+\\s*$") %>%    # trailing number at end of string
  as.numeric() %>%
  { .[!is.na(.) & . > 0] }

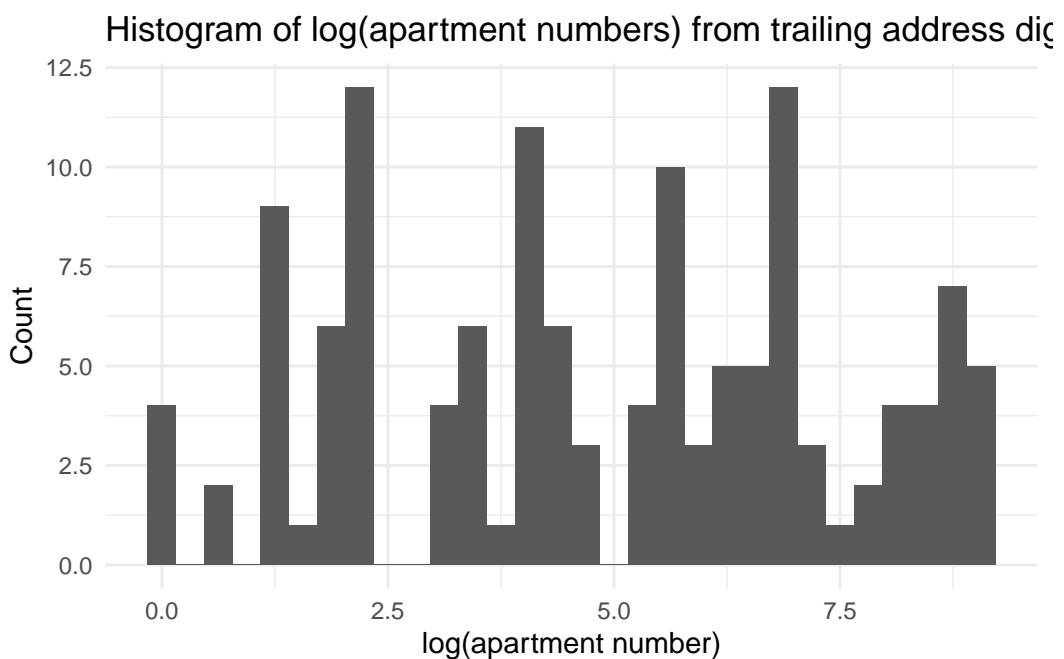
# Log-transform (natural log)
```

```

log_apt <- log(apt_num)

# Plot histogram (ggplot2)
ggplot(data.frame(log_apt = log_apt), aes(x = log_apt)) +
  geom_histogram(bins = 30, linewidth = 0.2) +
  labs(
    title = "Histogram of log(apartment numbers) from trailing address digits",
    x = "log(apartment number)",
    y = "Count"
  ) +
  theme_minimal()

```



## Part f

```

# Leading digit via logs (no string ops needed)
ld <- floor(apt_num / (10 ^ floor(log10(apt_num))))
ld <- ld[ld %in% 1:9]

obs_counts <- table(factor(ld, levels = 1:9))
n           <- sum(obs_counts)
obs_prop   <- as.numeric(obs_counts) / n

```

```

exp_prop <- log10(1 + 1 / (1:9))

# Chi-squared goodness-of-fit vs Benford
chisq.test(as.numeric(obs_counts), p = exp_prop, rescale.p = TRUE)

```

```

Chi-squared test for given probabilities

data: as.numeric(obs_counts)
X-squared = 62.268, df = 8, p-value = 1.67e-10

```

```

# (optional) quick summary
print(data.frame(
  digit = 1:9,
  observed = as.numeric(obs_counts),
  expected = round(n * exp_prop, 1),
  obs_prop = round(obs_prop, 4),
  exp_prop = round(exp_prop, 4)
))

```

	digit	observed	expected	obs_prop	exp_prop
1	1	12	39.1	0.0923	0.3010
2	2	19	22.9	0.1462	0.1761
3	3	16	16.2	0.1231	0.1249
4	4	13	12.6	0.1000	0.0969
5	5	14	10.3	0.1077	0.0792
6	6	18	8.7	0.1385	0.0669
7	7	8	7.5	0.0615	0.0580
8	8	11	6.6	0.0846	0.0512
9	9	19	5.9	0.1462	0.0458

```

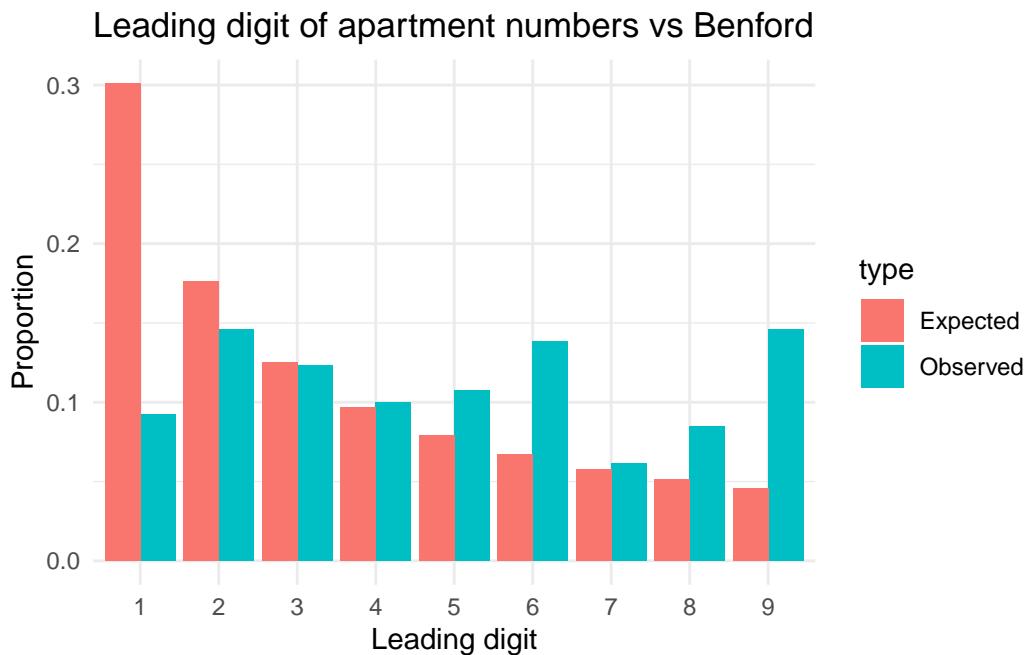
# (optional) simple plot
library(ggplot2)
library(tidyr)

plot_df <- data.frame(digit = factor(1:9), Observed = obs_prop, Expected = exp_prop) |>
  pivot_longer(c(Observed, Expected), names_to = "type", values_to = "prop")

ggplot(plot_df, aes(digit, prop, fill = type)) +
  geom_col(position = "dodge") +

```

```
labs(title = "Leading digit of apartment numbers vs Benford",
     x = "Leading digit", y = "Proportion") +
theme_minimal()
```



The apartment numbers do not follow Benford's law. The leading digit “1” occurs far less than Benford's 30% expectation, while higher digits (especially 6–9) occur more often, and the distribution lacks the characteristic monotonic decline. These systematic deviations indicate the data would likely fail a Benford goodness-of-fit test. This is consistent with apartment numbers being administratively assigned sequences rather than naturally generated measurements.