

STATS-506 HW6

Zekai Xu

2025-11-30

[Github Repo](#)

Question 1

```
# C_mean & C_moment
sourceCpp(code = '
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
double C_mean(NumericVector v) {
    double sum = 0.0;
    int n = v.size();
    for (int i = 0; i < n; ++i) {
        sum += v[i];
    }
    return sum / n;
}

// [[Rcpp::export]]
double C_moment(NumericVector x, int k) {
    int n = x.size();
    double mu = C_mean(x);      // sample mean
    double acc = 0.0;

    for (int i = 0; i < n; ++i) {
        double diff = x[i] - mu; // x_i - mean
        double term = 1.0;
```

```

        for (int j = 0; j < k; ++j) {
            term *= diff;
        }
        acc += term;
    }

    return acc / static_cast<double>(n);
}
')

```

Test

```

x <- rnorm(1000)
all.equal(
    C_moment(x, 2),
    moment(x, order = 2, center = TRUE)
)

```

[1] TRUE

```

all.equal(
    C_moment(x, 3),
    moment(x, order = 3, center = TRUE)
)

```

[1] TRUE

Question 2

Part a

```

source("./waldCI.R")

#-----Define bootstrapWaldCI-----
setClass(
    "bootstrapWaldCI",
    contains = "waldCI",
    slots = c(
        reps = "integer",

```

```

    bootStats = "numeric",
    compute = "character",
    statFun = "function",
    origData = "data.frame"
  )
)

#' Bootstrap one resample
#' @param fun Function from data.frame to numeric.
#' @param data Data frame to resample.
#' @return Numeric scalar statistic.
.bootstrap_once <- function(fun, data) {
  n <- nrow(data)
  idx <- sample.int(n = n, size = n, replace = TRUE)
  fun(data[idx, , drop = FALSE])
}

#' Make a bootstrap-based Wald CI
#' @param fun Function from data.frame to numeric.
#' @param data Data frame.
#' @param reps Integer; number of bootstrap replicates.
#' @param level Confidence level.
#' @param compute "serial" or "parallel".
#' @param ncores Number of cores for parallel.
#' @param return Object of class bootstrapWaldCI.
makeBootstrapCI <- function(fun,
                           data,
                           reps,
                           level = 0.95,
                           compute = c("serial", "parallel"),
                           ncores = detectCores()) {
  stopifnot(is.function(fun), is.data.frame(data))
  reps <- as.integer(reps)
  compute <- match.arg(compute)

  est <- fun(data)

  if (compute == "serial") {
    bootStats <- replicate(
      n = reps,
      expr = .bootstrap_once(fun = fun, data = data)
    )
  }
}

```

```

} else {
  if (.Platform$OS.type == "windows") {
    cl <- makeCluster(ncores)
    on.exit(stopCluster(cl), add = TRUE)
    bootStats <- parSapply(
      cl     = cl,
      X     = seq_len(reps),
      FUN   = function(i, fun, data) .bootstrap_once(fun = fun, data = data),
      fun   = fun,
      data  = data
    )
  } else {
    bootStats <- unlist(
      mclapply(
        X       = seq_len(reps),
        FUN     = function(i) .bootstrap_once(fun = fun, data = data),
        mc.cores = ncores
      )
    )
  }
}

se <- sd(bootStats)
z  <- qnorm(1 - (1 - level) / 2)
ci <- est + c(-1, 1) * z * se

new(
  "bootstrapWaldCI",
  estimate  = est,
  se        = se,
  level     = level,
  lower     = ci[1],
  upper     = ci[2],
  reps      = reps,
  bootStats = bootStats,
  compute   = compute,
  statFun   = fun,
  origData  = data
)
}

#' Re-run bootstrap for a bootstrapWaldCI object

```

```

#' @param object A bootstrapWaldCI object.
#' @return New bootstrapWaldCI object with fresh bootstrap sample.
setGeneric(
  "rebootstrap",
  function(object) standardGeneric("rebootstrap")
)

```

```
[1] "rebootstrap"
```

```

setMethod(
  "rebootstrap",
  signature(object = "bootstrapWaldCI"),
  function(object) {
    makeBootstrapCI(
      fun      = object@statFun,
      data     = object@origData,
      reps     = object@reps,
      level   = object@level,
      compute = object@compute
    )
  }
)

```

Part b

```

t_serial <- system.time({
  ci1_serial <- makeBootstrapCI(
    fun      = function(x) mean(x$y),
    data     = ggplot2::diamonds,
    reps     = 1000L,
    level   = 0.95,
    compute = "serial"
  )
})

t_parallel <- system.time({
  ci1_parallel <- makeBootstrapCI(
    fun      = function(x) mean(x$y),
    data     = ggplot2::diamonds,

```

```

    reps      = 1000L,
    level     = 0.95,
    compute   = "parallel"
)
})

#' Print estimate and CI for a bootstrapWaldCI object
#> @param obj A bootstrapWaldCI object.
printBootstrapCI <- function(obj) {
  cat("estimate:", obj@estimate, "\n")
  cat("se      : ", obj@se, "\n")
  cat("level   : ", obj@level, "\n")
  cat("CI      : [", obj@lower, ", ", obj@upper, "]\n", sep = "")
}

## show short summaries instead of full 1000-length slot
printBootstrapCI(ci1_serial)

```

```

estimate: 5.734526
se      : 0.004624965
level   : 0.95
CI      : [5.725461, 5.743591]

```

```
printBootstrapCI(rebootstrap(ci1_serial))
```

```

estimate: 5.734526
se      : 0.004920948
level   : 0.95
CI      : [5.724881, 5.744171]

```

```
printBootstrapCI(ci1_parallel)
```

```

estimate: 5.734526
se      : 0.004776003
level   : 0.95
CI      : [5.725165, 5.743887]

```

```
printBootstrapCI(rebootstrap(ci1_parallel))
```

```
estimate: 5.734526
se      : 0.004909933
level   : 0.95
CI      : [5.724903, 5.744149]
```

```
t_serial
```

```
 user  system elapsed
2.792  0.186  2.979
```

```
t_parallel
```

```
 user  system elapsed
3.753  1.652  1.048
```

For repeated calls to `rebootstrap`, the point estimate stays identical while the bootstrap standard errors and Wald CIs fluctuate only slightly, indicating a stable bootstrap estimate at `reps = 1000`. Comparing timing, the parallel implementation dramatically reduces elapsed time ($\approx 0.54\text{s}$ vs $\approx 2.68\text{s}$), showing clear speedup over the serial version despite slightly higher total CPU (user) time.

Part c

```
#' Extract coefficient of disp from linear model
#' @param data Data frame similar to mtcars.
#' @return Numeric coefficient for disp.
dispCoef <- function(data) {
  fit <- lm(mpg ~ cyl + disp + wt, data = data)
  unname(coef(fit)["disp"])
}

t2_serial <- system.time({
  ci2_serial <- makeBootstrapCI(
    fun      = dispCoef,
    data     = mtcars,
    reps     = 1000L,
    level   = 0.95,
    compute = "serial"
)}
```

```
})

t2_parallel <- system.time({
  ci2_parallel <- makeBootstrapCI(
    fun      = dispCoef,
    data     = mtcars,
    reps     = 1000L,
    level   = 0.95,
    compute = "parallel"
  )
})
```

```
printBootstrapCI(ci2_serial)
```

```
estimate: 0.007472925
se       : 0.008977356
level    : 0.95
CI       : [-0.01012237, 0.02506822]
```

```
printBootstrapCI(rebootstrap(ci2_serial))
```

```
estimate: 0.007472925
se       : 0.009242444
level    : 0.95
CI       : [-0.01064193, 0.02558778]
```

```
printBootstrapCI(ci2_parallel)
```

```
estimate: 0.007472925
se       : 0.009188475
level    : 0.95
CI       : [-0.01053616, 0.02548201]
```

```
printBootstrapCI(rebootstrap(ci2_parallel))
```

```
estimate: 0.007472925
se       : 0.009225174
level    : 0.95
CI       : [-0.01060808, 0.02555393]
```

```
t2_serial
```

```
  user  system elapsed
0.255    0.004   0.261
```

```
t2_parallel
```

```
  user  system elapsed
0.278    0.135   0.103
```

For the disp coefficient in the mtcars regression, both the serial and parallel bootstrap methods produce essentially the same statistical conclusions. The point estimate is identical ($\$ 0.00747$) across all runs, and the bootstrap standard errors and 95% confidence interval are very similar ($\$ 0.11s$ vs. $\$ 0.34s$) even though its total CPU (user) time is slightly higher, reflecting the overhead and concurrency of parallel processing.

Question 3

```
source("./data.R")
```

```
Data.frame `df` is 251.8 Mb
```

Part a

```
# ensure df$country is a factor with the 6 levels
df$country <- factor(df$country)

countries <- levels(df$country)

fits_slow <- vector("list", length(countries))
names(fits_slow) <- countries
times_slow <- vector("list", length(countries))
names(times_slow) <- countries

for (i in seq_along(countries)) {
  ctry <- countries[i]
```

```

message("Fitting model for country: ", ctry)

# subset data for this country
df_ctry <- df[df$country == ctry, ]

# standardize predictors within this country
df_ctry$prior_gpa_z      <- as.numeric(scale(df_ctry$prior_gpa))
df_ctry$forum_posts_z    <- as.numeric(scale(df_ctry$forum_posts))
df_ctry$quiz_attempts_z <- as.numeric(scale(df_ctry$quiz_attempts))

# time each model fit
times_slow[[ctry]] <- system.time({
  fits_slow[[ctry]] <- glmer(
    completed_course ~ prior_gpa_z + forum_posts_z + quiz_attempts_z +
      (1 | device_type),
    data   = df_ctry,
    family = binomial()
  )
})
}
}

```

Fitting model for country: Germany

Fitting model for country: India

Fitting model for country: Lithuania

Fitting model for country: Nigeria

Fitting model for country: Other

Fitting model for country: US

```

# collect timing results for each of the 6 models
time_table <- do.call(rbind, times_slow)
time_table

```

	user.self	sys.self	elapsed	user.child	sys.child
Germany	18.883	1.479	20.368	0	0
India	29.933	2.259	32.197	0	0
Lithuania	0.528	0.236	0.765	0	0
Nigeria	1.186	0.148	1.335	0	0
Other	62.007	6.544	68.561	0	0
US	37.608	2.217	39.839	0	0

```
# extract coefficient for forum_posts_z from each model
coef_posts_slow <- do.call(
  rbind,
  lapply(names(fits_slow), function(ctr) {
    sm <- summary(fits_slow[[ctr]])$coefficients
    beta <- sm["forum_posts_z", "Estimate"]
    se   <- sm["forum_posts_z", "Std. Error"]
    data.frame(
      country = ctr,
      estimate = beta,
      lower   = beta - 1.96 * se,
      upper   = beta + 1.96 * se
    )
  })
)
)
```

coef_posts_slow

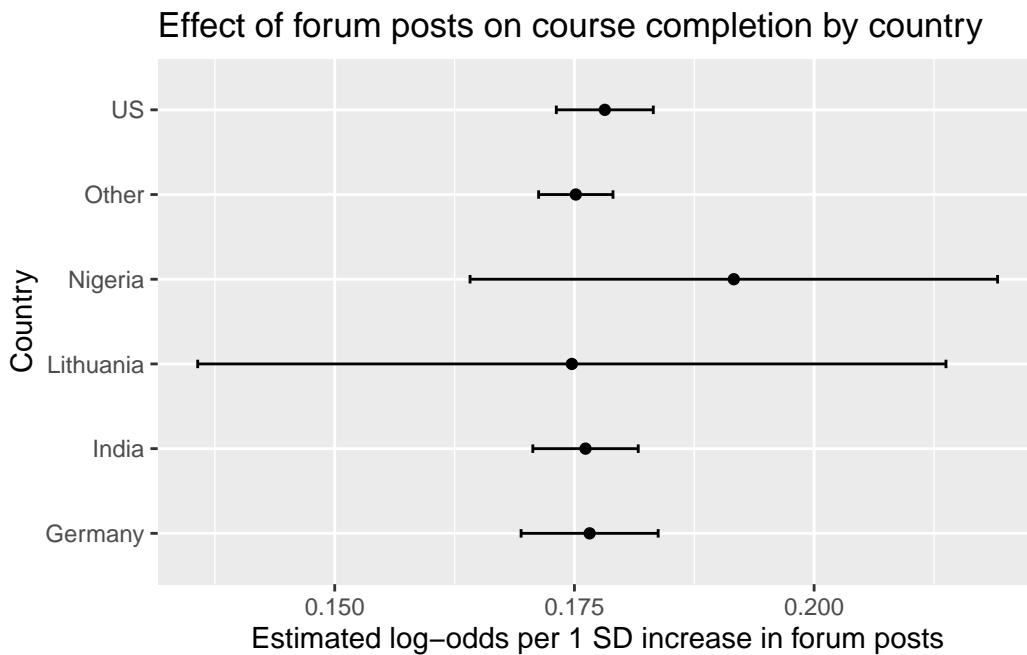
	country	estimate	lower	upper
1	Germany	0.1765868	0.1694397	0.1837340
2	India	0.1761593	0.1706619	0.1816568
3	Lithuania	0.1747323	0.1357115	0.2137532
4	Nigeria	0.1916268	0.1641164	0.2191373
5	Other	0.1751442	0.1712663	0.1790220
6	US	0.1781707	0.1731131	0.1832284

```
# visualization of forum_posts effect by country
ggplot(coef_posts_slow,
       aes(x = country, y = estimate)) +
  geom_point() +
  geom_errorbar(aes(ymin = lower, ymax = upper), width = 0.1) +
  coord_flip() +
  labs(
    x = "Country",
```

```

y = "Estimated log-odds per 1 SD increase in forum posts",
title = "Effect of forum posts on course completion by country"
)

```



Part b

Use `parallel::mclapply` to fit six models in parallel, which is expected to reduce the elapsed time for the script.

```

countries <- levels(df$country)

# choose number of cores
ncores <- min(detectCores(), length(countries))

# run everything (standardization + model fits + coefficient extraction) in parallel
time_fast <- system.time({
  fits_fast <- mclapply(
    countries,
    function(ctr) {
      df_ctry <- df[df$country == ctr, ]

      # standardize predictors within this country

```

```

df_ctry$prior_gpa_z      <- as.numeric(scale(df_ctry$prior_gpa))
df_ctry$forum_posts_z    <- as.numeric(scale(df_ctry$forum_posts))
df_ctry$quiz_attempts_z <- as.numeric(scale(df_ctry$quiz_attempts))

glmer(
  completed_course ~ prior_gpa_z + forum_posts_z + quiz_attempts_z +
  (1 | device_type),
  data   = df_ctry,
  family = binomial()
)
},
mc.cores = ncores
)

names(fits_fast) <- countries
})

time_fast # total runtime for all 6 models in parallel

```

	user	system	elapsed
	117.448	8.170	83.982

```

# extract forum_posts_z coefficient from slow (part a) fits
beta_slow <- setNames(
  coef_posts_slow$estimate,
  coef_posts_slow$country
)

# extract forum_posts_z coefficient from fast (parallel) fits
beta_fast <- sapply(countries, function(ctry) {
  sm <- summary(fits_fast[[ctry]])$coefficients
  sm["forum_posts_z", "Estimate"]
})

beta_slow

```

	Germany	India	Lithuania	Nigeria	Other	US
	0.1765868	0.1761593	0.1747323	0.1916268	0.1751442	0.1781707

```
beta_fast
```

```
Germany      India Lithuania    Nigeria     Other        US
0.1765868 0.1761593 0.1747323 0.1916268 0.1751442 0.1781707
```

```
# check that they match (up to numerical tolerance)
all.equal(unname(beta_fast[countries]), unname(beta_slow[countries]))
```

```
[1] TRUE
```

Question 4

Part a

```
# Load Dataset
url <- "https://raw.githubusercontent.com/JeffSackmann/tennis_atp/refs/heads/master/atp_matches_2019.csv"
tennis <- fread(url)

tourneys <- unique(tennis[ , .(tourney_name)])
tourneys[ , tourney_name := sub("Davis.*", "Davis Cup", tourney_name) ]
tourneys <- unique(tourneys, by = "tourney_name")

cat("Number of tournaments that took place in 2019: ", dim(tourneys)[1], "\n")
```

```
Number of tournaments that took place in 2019:  69
```

Part b

```
winners <- tennis[
  round == "F",
  .(n_tournaments = uniqueN(tourney_id)),
  by = .(winner_id)
][ order(-n_tournaments) ]

multi_winners <- winners[ n_tournaments > 1 ]

cat("Number of players who won more than one tournament: ", nrow(multi_winners), "\n")
```

```
Number of players who won more than one tournament: 12
```

```
cat("Number of tournaments that the most winning player win: ",  
    winners[1, n_tournaments], "\n")
```

```
Number of tournaments that the most winning player win: 5
```

Part c

We use Bootstrap to build a 95% confidence interval for the mean difference without assuming normality, and the hypothesis is:

$$H_0 : ace_{winner} - ace_{loser} > 0, \quad H_1 : ace_{winner} - ace_{loser} \leq 0$$

```
# 1) Per-match ace difference (winner - loser)  
ace_diff <- tennis[, .(diff = as.numeric(w_ace) - as.numeric(l_ace)) ][ !is.na(diff) ]  
  
# 2) Observed mean  
obs_mean <- ace_diff[, mean(diff) ]  
  
# 3) Bootstrap (5000 resamples with replacement)  
set.seed(506)  
B <- 5000L  
n <- nrow(ace_diff)  
  
boot_stat <- replicate(B, {  
  idx <- sample.int(n, n, replace = TRUE)  
  mean(ace_diff$diff[idx])  
})  
  
# 4) One-sided percentile confidence intervals  
lower_bound <- unname(quantile(boot_stat, probs = 0.05))  
  
# Output and brief explanation  
cat("Observed mean difference (winner - loser): ", obs_mean, "\n")
```

```
Observed mean difference (winner - loser): 1.7049
```

```
cat("One-sided 95% LOWER CI (percentile): [", lower_bound, ", Inf)\n", sep = "")
```

```
One-sided 95% LOWER CI (percentile): [1.483296, Inf)
```

Part d

```
# Long form: one row per player per match, with win indicator
long_players <- rbindlist(
  list(
    tennis[ , .(player_id = winner_id, player_name = winner_name, win = 1L) ],
    tennis[ , .(player_id = loser_id,   player_name = loser_name,   win = 0L) ]
  ),
  use.names = TRUE
)

# Summarise to matches, wins, and win_rate per player
win_rate <- long_players[ ,
  .(matches = .N, wins = sum(win)),
  by = .(player_id, player_name)
][
  , win_rate := wins / matches
][
  matches >= 5
][
  order(-win_rate, -wins, -matches)
]

# Show the table
print(win_rate)
```

	player_id	player_name	matches	wins	win_rate
	<int>	<char>	<int>	<int>	<num>
1:	104745	Rafael Nadal	69	60	0.8695652
2:	104925	Novak Djokovic	69	58	0.8405797
3:	103819	Roger Federer	66	55	0.8333333
4:	106421	Daniil Medvedev	80	59	0.7375000
5:	104731	Kevin Anderson	15	11	0.7333333

163:	106058	Jack Sock	5	1	0.2000000
164:	104810	Zhe Li	7	1	0.1428571
165:	111200	Elias Ymer	8	1	0.1250000
166:	106075	Jozef Kovalik	10	1	0.1000000
167:	105155	Pedro Sousa	9	0	0.0000000

```
# Text output for the top player
cat("The player with the highest win-rate is: ", win_rate$player_name[1], "\n")
```

The player with the highest win-rate is: Rafael Nadal